



Tutoring Section 13

Machine Learning: Correlation, Regression

Slides by Kevin Miao

Logistics

- Vibe Check:
 - How stressed/relaxed do you feel?
 - 2.5 weeks left of classes! How prepared do you feel for the last stretch of the class/semester?
- **Project 3**, movie recommendations, has been released!

As always, let me know if you have any questions about anything.

Today

- Correlation Coefficient
 - Regression
 - Errors: (R)MSE
-


Worksheet

Link: <https://tinyurl.com/d8tutweek13>

Associations

- **Association**

- Any type of relationship between two variables
 - Could be **linear**, **non-linear**



- In this class, we will only focus on **linear relationships**
-

Correlation

- **Correlation**

- **Goal:** How do we quantify a linear relationship?

- **Correlation coefficient, r**

- Strength
 - Direction

- **Calculation**

- **Mean of the product of x and y in standard units**

- Does our correlation coefficient change if ...

- We swap our axes x , y ?

- We convert our x units from say inches to centimeters?

- What is the range of our correlation coefficient?
-

Q1.1a

Practice Problems

1.1 The following table, `taters`, depicts the number of tater tots a person has eaten, along with a number that quantifies their satisfaction, which is a number that goes from 0 to 10.

Tater Tots Consumed	Satisfaction
1	8
10	3
4	7
3	10
7	6
3	8

a) Complete the function `standard_units` which takes in an array `num_array` and returns the same array in standard units.

```
def standard_units(num_array):  
    arr_mean = _____  
    arr_sd = _____  
    return _____
```

Q1.1bcd

b) Fill in the blanks to define a function `correlation` that finds the correlation from a table. It takes in three arguments: a table, `tbl`, and two column indices, `x` and `y`.

Hint: Use the `standard_units` function defined above!

```
def correlation(tbl, x, y):  
    su_x = _____ ( _____ )  
    su_y = _____ ( _____ )  
    return _____ ( _____ )
```

c) Calculate `r` by using the `correlation` function.

```
correlation(_____, _____, _____)
```

d) Suppose that we calculated a value of r to be equal to -0.879. What can you conclude about the association between the number of tater tots consumed and a person's satisfaction?

Q1.2

1.2 True or False?


- a. A high value of r shows that a change in x causes a change in y .

 - b. If we switch the axes of a plot, the correlation coefficient will not change.

 - c. Suppose that we calculated a value of r to be equal to .83. We should conclude that eating taters is indeed correlated with satisfaction.
-

Q1.3

1.3 Answer the following questions about the plots below.



- a. Order the scatter plots above in from least correlated to most correlated.
- b. Which plots have a positive correlation coefficient? Negative correlation coefficient?

Regression

- **Objective:** We want to predict a **number** based on given parameters.
 - **Linear Regression**
 - We know that the relationship between our variables and the number we want to predict has a linear shape.
 - **Calculating the formula/line that predicts the numbers**
 - Calculate the correlation coefficient
 - Mean of the product of x and y in standard units
 - Calculate the slope
 - $\text{Slope} = r * (\text{SD_Y}/\text{SD_X})$
 - Calculate the intercept by plugging in the means of x and y
-

Q2.2


Practice Problems

The `water` table contains one row per country with data from 2014. The `OBS_VALUE` column represents the approximate price ranking of a 1.5 liter bottle of mineral water in that country, and the `mm_precipitation` column represents the average precipitation in that country (in millimeters).

COUNTRY	OBS_VALUE	mm_precipitation
Albania	0.55	1485
Algeria	0.27	80

... (89 rows omitted)

Expression	Values
<code>np.average(water.column('OBS_VALUE'))</code>	0.919016
<code>np.std(water.column('OBS_VALUE'))</code>	0.464763
<code>np.average(water.column('mm_precipitation'))</code>	1010.4
<code>np.std(water.column('mm_precipitation'))</code>	752.475
<code>correlation(water, 'OBS_VALUE', 'mm_precipitation')</code>	0.262079



2.2 Write an equation for the regression line of the data in the `water` table, using `OBS_VALUE` as y using the `mm_precipitation` as x .

Q2.3


Practice Problems

The `water` table contains one row per country with data from 2014. The `OBS_VALUE` column represents the approximate price ranking of a 1.5 liter bottle of mineral water in that country, and the `mm_precipitation` column represents the average precipitation in that country (in millimeters).

COUNTRY	OBS_VALUE	mm_precipitation
Albania	0.55	1485
Algeria	0.27	80

... (89 rows omitted)

Expression	Values
<code>np.average(water.column('OBS_VALUE'))</code>	0.919016
<code>np.std(water.column('OBS_VALUE'))</code>	0.464763
<code>np.average(water.column('mm_precipitation'))</code>	1010.4
<code>np.std(water.column('mm_precipitation'))</code>	752.475
<code>correlation(water, 'OBS_VALUE', 'mm_precipitation')</code>	0.262079



2.3 Using the regression line equation above, what would we expect the `OBS_VALUE` to be in 2014 for a country that had an average of 700 mm of precipitation?

Errors

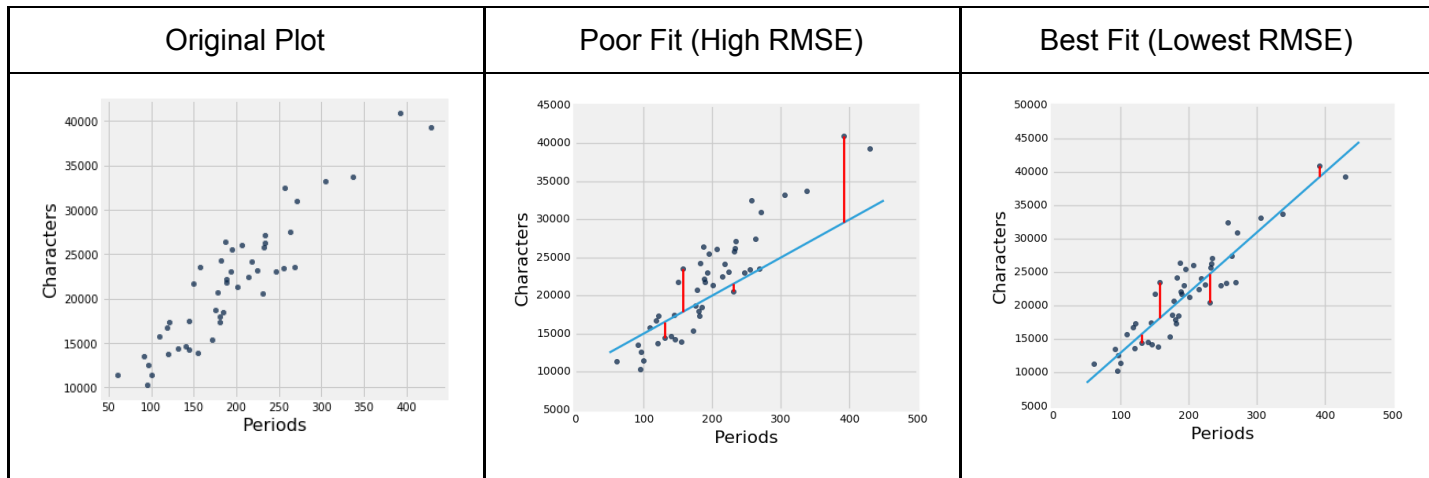
- **Context:** In Data 8, we provide you a lot of statistical knowledge, but in traditional Machine Learning Engineering side, we will approach ML problems through this perspective.
 - **Set-up:** You have a problem, you want to predict something, define a model (Linear Regression), define an error (RMSE/MSE) and minimize it. This gives you a model (line) with the lowest error possible given your data points.
 - **Here:**
 - **Root Mean Squared Errors** (looks like SD)
 - **Square root of the average of squared errors**
 - $\sqrt{(\text{error}_{point\ 1}^2 + \dots + \text{error}_{point\ n}^2)/n}$
 - Error = actual y – predicted y
-

RMSE/MSE/Linear Regression Facts

- **Statistics perspective vs Computer Science perspective**
 - The line calculated with correlation coefficients is the same line that minimizes the error. In other words, the linear regression line is the line that is the best!
- **Why do we pick RMSE?**
 - It is completely in your right to substitute the RMSE by another loss function such as the absolute loss. It provides different assumptions.
- **What does the minimize function do?**
 - Imagine it takes the derivate, sets it to 0 and calculates the parameters for which the maximum is attained.
- **What happens if we run minimize on MSE instead of RMSE?**
 - MSE does not change the shape of the graph and will not affect the outputted line.

You don't have to know what's in grey.

Examples



Q3.1-3.2

Practice Problems

3.1 Write a function that returns the RMSE of an array of observed values if the predicted values are given by an array. The two arrays have the same length.

```
def RMSE(observed, predicted):  
    residual = _____  
    squared_residuals = _____  
    squared_resid_avg = _____  
    return _____
```

3.2 In the calculation of root mean squared error, why is it important for us to square the residual before taking the sum?
